

MEETING REPORT

Date: May18th, 2026

Location: Remote Conference (Zoom)

Attendees:

- The Client (Delivery Company CEO)
- The Engineering Team (Kevin Otah Ogbusuo, Mishvaba Hitensinh Vaghela, Pritheesh Selvaraja Kumar)

1. Client Interview Full Transcript

2. Meeting Summary

3. Follow-up Actions

CLIENT INTERVIEW FULL TRANSCRIPT

Engineering team:

First of all, we'd like to get general overview of your company. What are the pain points exactly. We see that expansion is increasing complexity of delivery organisation, but how was the company handling it in the past and what exactly is changing? All the constraints mentioned should have always been there to begin with, how did the company handle them? With the growth of the company, were there new constraints/limitations? Basically, what exactly is making it "increasingly complex"?

Client:

Sure. So, at the beginning, honestly, things were simple because the volume was low and everything fit in my head. We had fewer clients, mostly pharmacies, no real time windows, and if something went wrong I could just reshuffle a route on the fly or jump in the van myself. Excel worked because there weren't that many deliveries, and the drivers all knew the area by heart.

What's changing now is not one new constraint, it's the **combination** of them happening at the same time. We still have the same fleet size, roughly the same number of drivers, but more stops per day, more medical urgency, and more clients with specific expectations. The hospital before 9am, temperature-sensitive products, traffic in Paris, a driver calling in sick on a Monday — each of those existed before, but not all together, every week.

In the past, when we had a delay, it affected maybe one client and we could smooth it over. Now, one small disruption cascades: if the refrigerated vehicle is late, it's not interchangeable with another van; if a morning delivery slips, it can break the whole afternoon route. Planning used to be "good enough"; now it's fragile. That's what feels increasingly complex.

So yes, the constraints were always there in theory, but growth removed our margin for improvisation. We're still managing, but honestly it's a mess some days — especially Mondays — and I spend more time firefighting than actually running operations.

Engineering team:

What does just increasing fleet size and number of drivers do to help your situation? For instance, if the refrigerated vehicle is late, use another, and if a driver calls in sick, another takes his place. Would this help or is it just using band-aid on a broken arm? If no, to what extent would this contribute? And, what exactly about Monday makes it a mess?

Client:

Increasing the fleet does help a bit, but it's mostly a band-aid. If I had a second refrigerated vehicle, yes, a late one wouldn't block everything — but that vehicle would sit idle most of the week, and it's expensive to run. Same with drivers: having a "spare" driver sounds great on paper, but with our volumes, that person would be under-used four days out of five. So it reduces stress, but it doesn't really fix the root problem.

The deeper issue is coordination, not just capacity. Even with more vans, you still need to decide who goes where, in what order, and with which constraints. When a driver calls in sick, it's not just "swap another one in" — the replacement may not know that hospital route, or the refrigerated constraints, or the shortcuts in Paris traffic. So adding resources helps maybe 20–30%, but complexity still grows faster than headcount.

Mondays are a mess because everything peaks at once. Volume is higher after the weekend, the hospital always has early deliveries, pharmacies are restocking, and that's also the day drivers are most likely to call in sick. Planning that morning feels like juggling glass balls — one small issue at 7:30am and the whole day is off balance. Honestly, that's when Excel really starts to break down.

Engineering team:

Ah, so this is more a management/scheduling problem than human/fleet problem. Basically you want to be able to use to the best capacity what you already have, right? And in such way that you can move things around, have a few setbacks and things still run okay at least. You need a system, something you can "plug" values into and it, based on ALL the data/information you have, gives a good schedule to use or shuffle around based on any other additional data. Are we in the ballpark? If no, how does that differ from what you might have in mind? But before we get there, we have to understand all aspects of the situation. So about the decision of who, what, where, etc... how's this: What if addresses are sorted based on distance and consequently based on the time based on the time it takes to get there? This provides an idea on the order of deliveries that is the most efficient. If we know the time it takes to make deliveries, we should be able to allocate them more confidently to drivers based on their schedule/time constraints. How does this sound? (For now, we are not considering some potential issues like traffic, vehicle breakdown for now. For the sake of finding at least a bare bones solution) What exactly about Excel doesn't handle your needs appropriately? What can/needs to be improved? What is your ideal?

Client:

Yes, you're very much in the ballpark. What I'm looking for is exactly that: being able to use what we already have better, and not have everything collapse the moment there's a small setback. I don't expect perfection, but I want something more resilient than "Sophie panics at 7:45 and rewrites half the plan."

On your idea of sorting addresses by distance and time: in theory, that makes total sense, and that's roughly what I try to do in my head. The problem is that "most efficient order" on paper often clashes with reality — one stop has a strict time window, another needs the refrigerated van, another driver refuses central Paris after 10am because of traffic. So distance is a good base, but it's never the only factor, and Excel doesn't help me juggle all of that at once.

That's really where Excel breaks down. It's fine for listing deliveries, drivers, and vehicles, but the moment I change one thing — a sick driver, an urgent hospital delivery — I'm manually moving rows, recalculating times, double-checking constraints. There's no warning saying "this breaks a time window" or "this driver is now over hours." It relies entirely on my attention, and on busy days, that's risky.

My ideal? Something that gives me a solid first plan in the morning, and that I can adjust without everything unraveling. I still want to make decisions — I trust my drivers — but I want fewer blind spots and less mental load. Right now, honestly, the system is me, my Excel, and a lot of stress.

Engineering team:

Alright, so basically, there are several other factors in play. So before we suggest any other ideas, let's talk about them: HOW do you decide what deliveries are the most important or urgent? Are there clear indicators or is it just based on the situation and your intuition/experience? What do you usually consider when planning deliveries? We're guessing distance and time are just part of it. What are the usual causes of the "mess" you speak of? Which are the biggest? About the 24 hour time limit, can a destination actually be 24 hours away? If the stock is loaded onto the vehicle, is it imperative that it gets to the destination before it can be offloaded? Basically, are there drop-offs or corresponding spots? Or branches? Is there something we're missing/overlooking so far? Also, give me some scale/numbers from your business. How many deliveries do you usually handle? How does that number differ from a chaotic day like Mondays? For your drivers and fleet, are they usually available? As in, how many? Are refrigerated vehicles usually an issue? How many do you have?

Client:

Yes, you're right to slow down and really map this out first. Let me answer point by point, from how we decide urgency to the scale of operations.

Urgency / importance Some deliveries are clearly urgent, others less so. The hospital is the clearest case: they have a strict time window, usually before 9am, and if we miss it, I get a call immediately. Temperature-sensitive products are also high priority — once they leave the warehouse in the refrigerated vehicle, I don't want them sitting around. For the rest, pharmacies are more flexible, but even then it's a relationship thing: some expect morning delivery, others don't care as long as it's same day. We don't have a formal "priority score"; it's mostly experience and habit.

What I consider when planning Distance and time are the starting point, yes, but quickly other things take over:

- time windows (about 30% of deliveries)
- vehicle type (refrigerated vs normal)
- driver familiarity with certain routes or clients
- legal working hours — I really try not to push people to the limit
- and honestly, client sensitivity: who will complain first if we're late

All of that lives in my head or scattered in Excel notes.

What usually causes the “mess” The biggest causes are:

- a driver calling in sick in the morning
- one urgent delivery being added last minute
- the refrigerated vehicle running late and blocking everything else

When that happens, I'm not just moving one delivery, I'm reshuffling half the day. That's when mistakes happen.

About the 24-hour limit No, destinations are not 24 hours away — everything is Paris and inner suburbs. The 24 hours is a medical constraint: once the product leaves the warehouse, it must be delivered within 24 hours, no exceptions. We don't do intermediate drop-offs or storage. It's point-to-point: warehouse → client. So if it's loaded, it has to reach its destination.

Drop-offs / branches No branches, no relay points. Each delivery has a single destination. That's one thing that's actually simple.

Scale / numbers On a normal day, we handle around **40 to 60 deliveries**. Mondays are closer to the high end, sometimes even a bit above that if there's backlog from the weekend. Fridays are calmer.

We have **5 drivers**, each on 8-hour shifts, with a legal max of 10 hours if things go wrong. **6 vehicles total**:

- 3 small vans (good in Paris, but limited range and capacity)
- 2 large vans (better for suburbs, worse in traffic)
- 1 refrigerated vehicle

The refrigerated vehicle is a bottleneck. We only have one, and when several temperature-sensitive deliveries pile up, planning becomes very tight.

Are drivers and fleet usually available? Usually yes — but “usually” is doing a lot of work here. One absence already hurts because we don't have slack. Same for vehicles: there's no spare refrigerated van, and that's stressful.

From my side, what you're not overlooking is the human pressure. When things go wrong, it's not abstract — it's clients calling my phone, drivers waiting for instructions, and me trying not to make the wrong call. That's why I'm so focused on understanding all the constraints before even thinking about solutions.

Engineering team:

Alright, this paints a much clearer picture. We've got 5 drivers, 6 vehicles (with that one refrigerated bottleneck), and around 40-60 deliveries peaking on Mondays. The Excel situation makes total sense, it's just a static ledger, not a dynamic engine. It can't warn you when a decision breaks a rule, which is exactly why things cascade when there's a hiccup. Before we jump into designing this dynamic engine for you, We need to zoom out just a little bit to understand the broader environment you're operating in, and then zoom way into the actual numbers. First, about the business environment: Who are you actually competing against in this space? Are there giant logistics companies trying to eat your lunch, or is it mostly other specialized mid-sized players? And on the flip side, do you have specific partners like certain pharmaceutical suppliers, labs, or medical networks that sort of dictate how you operate? Second, let's talk about the tech and the data you're already using. Beyond your Excel sheet in the office, what else is in the technological ecosystem? How do the drivers actually navigate out there? Are they just using Waze or Google Maps on their personal phones, or do you have a dedicated system in the vans? And how do the delivery orders actually reach your desk in the first place? Do clients email you, or is there an order portal? Finally, to build a system that actually works mathematically, We need some specific numbers, the 'orders of magnitude' as we call them. When we talk about vehicle 'capacity', how do you actually measure that in your day-to-day? Are we talking about the number of standard medical boxes, weight, or volume? Roughly how much can a small van hold compared to a large one? Also, when a driver arrives at a hospital or pharmacy, how long does the actual drop-off take on average? Is it a quick 5-minute handoff or a 30-minute ordeal with paperwork? And what about loading up the vans at the warehouse in the morning?

Client:

Good questions — this is exactly the right level to be thinking at.

Competitive environment

We're not really competing with the big global logistics players; they're too rigid and too expensive for this kind of medical last-mile work. Our real competitors are other small or mid-sized specialized couriers who focus on pharma and clinics. What differentiates us today is reliability and relationships, not price — if we start missing time windows, that advantage disappears fast.

Partners / who dictates the rules

The constraints mostly come from the medical side: pharmaceutical suppliers, labs, and especially the hospital. The hospital's time window is non-negotiable, and suppliers impose the 24-hour rule on certain products. We don't have much leverage there; we adapt to them, not the other way around.

Current tech ecosystem

It's very basic. Orders come in by email or sometimes by phone — no portal. I consolidate everything manually into Excel in the morning. Drivers use their personal phones with Google Maps or Waze, and we coordinate changes via WhatsApp. There's nothing installed in the vans themselves.

Vehicle capacity — how we think about it

We don't measure capacity very scientifically. It's mostly "number of medical boxes" and experience. A small van can handle roughly 20–30 standard deliveries before it's full or inefficient, while a large van can take maybe 40–45 depending on box size. The refrigerated vehicle is similar to a large van in volume, but obviously restricted by product type.

Time on site (drop-offs)

Pharmacies are usually quick — 5 to 10 minutes if there's no queue. Clinics are similar. The hospital is the exception: that can easily be 20–30 minutes with waiting, signatures, and internal procedures. That's why missing the hospital window messes up an entire route.

Morning loading at the warehouse

Loading typically takes about 30–45 minutes in total, assuming everything is ready. If there's a last-minute order or something wasn't prepared properly, that's when we already start the day late — and once you're late in the morning, you're chasing the clock all day.

So overall, nothing here is insanely complex on its own — it's the accumulation. Right now, all of this lives in people's heads and in Excel cells, and that's exactly where it starts to crack under pressure.

Engineering team:

This is brilliant, thank you. The breakdown of drop-off times and capacities gives us exactly the kind of mathematical baseline we need. It's clear that the 'mess' isn't from a lack of resources, but rather the fragility of trying to balance all these intersecting rules in a static spreadsheet. We believe we have almost everything we need to start modeling a solution for you, but we just want to iron out the final details on the variables and your ultimate goals. First, regarding the actual driving time between stops in Paris and the suburbs—how much variability are we talking about? Do you currently use a rough average (like 15 minutes between stops) when planning in your head, or does the time of day completely change those estimates? Second, you mentioned earlier about 'driver familiarity' with certain routes or clients. Is this a strict rule (for example, Driver A must handle the central Paris route), or is it just a preference that makes the day go smoother? Finally, when our new system generates a route plan for you, what is the ultimate 'score' of a good day in your eyes? If we have to make a trade-off, what is your top priority? Is it finishing as early as possible to save driver hours, minimizing the total distance to save on fuel and vehicle wear, or purely hitting 100% of the time windows and keeping clients happy, regardless of the routing efficiency?

Client:

Yes, you're asking the right questions — these are exactly the trade-offs I'm making every day.

Driving time variability

Time of day changes everything. Early morning, Paris can be fine; after 9–10am, a

10-minute hop can easily become 25. In my head, I use very rough buckets: “early morning is fast”, “late morning is risky”, “afternoon in Paris is unpredictable”. I don’t use a fixed average like 15 minutes — it’s more instinct than math, which is part of the problem.

Driver familiarity

It’s not a strict rule, more a strong preference. Certain drivers know central Paris, parking spots, hospital procedures — putting someone else there doesn’t break the day, but it makes it slower and more stressful. On a calm day I respect that; on a bad Monday, I override it and accept some inefficiency.

What defines a ‘good day’

Top priority, no hesitation: **hitting time windows and avoiding angry client calls.** I’ll happily accept extra kilometers or a driver finishing a bit later if everyone gets their deliveries on time. Finishing early or saving fuel is nice, but secondary — a single missed hospital delivery costs me more than a week of fuel savings.

So if a system ever forces a trade-off, I’d always choose reliability and client trust over pure efficiency. Everything else is negotiable.

MEETING SUMMARY

The engineering team conducted a requirements-gathering interview with the CEO of a medium-sized medical delivery company operating in Paris and its suburbs. The company is experiencing scaling issues where the current manual planning system (Excel) is failing to handle complex, intersecting constraints (vehicle breakdowns, driver absences, strict time windows, and traffic variability). The CEO emphasized that the primary goal is not necessarily fuel efficiency, but rather absolute reliability and maintaining client trust by hitting strict delivery time windows.

FOLLOW-UP ACTIONS

1. Formalize the problem as a Vehicle Routing Problem with Time Windows (VRPTW).
2. Map the operational data (capacities, time on site, fleet availability) into machine-readable parameters.
3. Design a naive heuristic prototype to replace the manual Excel scheduling.
4. Draft a Product Requirements Document (PRD) detailing how to implement a deterministic optimization algorithm to scale the solution.